

# Logistic Regression

---

Doru Arfire


*Papers we love X, June 23rd 2016*

# Contents

- Linear Regression
- Logistic Regression
- Extensions to Logistic Regression
- Overfitting & Regularization
- Parameter selection
- Demo
- Q&A

# Ordinary Least Squares


- Given real values  $\mathbf{X}_{N \times K}$  and  $\mathbf{y}_N$ , train model that can predict  $\mathbf{y}'$  from  $\mathbf{X}'$
- $\mathbf{X}$  - independent variables
- $\mathbf{y}$  - dependent variable
- Assumes  $\mathbf{y}_i = \mathbf{w}^T \mathbf{x}_i + \mathbf{w}_0 + \epsilon_i$
- $\epsilon$  - error terms, assumed to have  $E[\epsilon] = 0$
- Columns (regressors) of  $\mathbf{X}$  linearly independent



# Ordinary Least Squares


- We're learning model given by  $\mathbf{f}(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$
- $\mathbf{x}_i = [\mathbf{1}, x_{i1}, x_{i2}, \dots, x_{iK}]^T$
- $\mathbf{w} = [\mathbf{w}_0, \mathbf{w}_1, \mathbf{w}_2, \dots, \mathbf{w}_K]^T$
- $\mathbf{w}_0$  is called the *intercept*
- We find  $\mathbf{w}$  by minimizing  $\|\mathbf{f}(\mathbf{x}) - \mathbf{y}\|^2$

$$w^* = \arg \max_w \sum_{i=1}^N (y_i - w^T x_i)^T (y_i - w^T x_i)$$




# OLS: beyond linearity

- Sometimes we suspect our input variables to have non-linear interactions
- E.g.,  $y$  depends not on  $x_1, x_2$ , etc., but on  $x_1^2, x_1x_2$ , etc.
- Take advantage of linear relationship between  $y$  and  $x_1^2, x_1x_2$
- Add extra variables  $z_1 = x_1^2, z_2 = x_1x_2$ , etc.




# Logistic Regression: Introduction

- Linear regression is good for, well, *regression*
- Unnatural fit for *classification*: given  $\mathbf{X}$ , predict  $\mathbf{y} \in \{0, 1\}$
- You could have a threshold  $\Theta$ ; anything above is 1, anything below is -1
- Difficult to interpret  $\mathbf{f}(\mathbf{x})$  as degree of certainty (what is  $f(x) = 10$ , or  $-200$ )
- How do we find  $\Theta$




# Generalized Linear Model



# Logistic Function

$$\text{sigm}_{\alpha}(x) = \frac{1}{1 + e^{-\alpha x}}$$


- Differentiable everywhere => we can do gradient descent
- Result is a probability-like value
- We can use it as a measure of certainty
- Parameterizable slope with  $\alpha$





# Logistic Function

- First discovered by Pierre François Verhulst (1845)
- Rediscovered by Raymond Pearl and Lowell Reed (1920s)
- Used to model population growth
- Introduced in statistical analysis by Berkson (*Application of the Logistic Function to Bio-Assay*, 1944)
- Initially, an alternative to **probit** - CDF of normal distribution



# Logistic regression: loss function


- Find  $\mathbf{W}$  by minimizing the number of errors  $\|y - \sigma(w^T x)\|$
- We interpret  $P(y = 1|x; w) = \sigma(w^T x)$
- Therefore  $P(y = 0|x; w) = 1 - \sigma(w^T x)$
- Equivalently  $P(y|x; w) = \sigma(w^T x)^y (1 - \sigma(w^T x))^{(1-y)}$
- We define the likelihood:  $L(w) = \prod_{x_i} P(y_i|x_i; w)$
- The cost function

$$J(w) = -\frac{1}{N} \sum_{i=1}^N y_i \log(\sigma(w^T x_i)) + (1 - y_i) \log(1 - \sigma(w^T x_i))$$

- Minimize it w.r.t.  $\mathbf{w}$  (MLE)


# Logistic Regression: discussion

- Bowl shaped error surface
- Gradient descent is guaranteed to find a global minimum
- Good results if data is linearly separable
- XOR problem, not linearly separable
- We need a more complex boundary



# Logistic regression: multinomial extensions

- How do we use LR for more than 2 classes?
- 1-vs-all approach:
- We can run  $K$  independent regressions
- Each will compute a different set of parameters
- Choose class with best result
- The probabilities need not sum to 1



# Logistic regression: Softmax

- A method to directly compute  $P(y = k | w; X)$ , using a single model
- Softmax function generalizes the logistic function to  $K$  classes

$$P(y_i = j | x_i; w) = \frac{e^{w_j^T x_i}}{\sum_{k=1}^K e^{w_k^T x_i}}$$

- $\mathbf{X}$  is  $N \times M$ ;  $\mathbf{w}$  is  $M \times K$
- The cost function becomes

$$J(w) = -\frac{1}{N} \sum_{i=1}^N \sum_{j=1}^K I\{y_i = j\} \log \frac{e^{w_j^T x_i}}{\sum_{l=1}^K e^{w_l^T x_i}}$$

# Logistic regression: Softmax

$$J(w) = -\frac{1}{N} \sum_{i=1}^N \sum_{j=1}^K I\{y_i = j\} \log \frac{e^{w_j^T x_i}}{\sum_{l=1}^K e^{w_l^T x_i}}$$

- Generalizes logistic regression cost function
- No closed form solution, solved using numeric optimization (GD)
- Still convex, GD will find a global maximum
- Overparameterized: multiple  $\mathbf{w}$  settings will optimize it
- Not equivalent to 1-vs-all approach
- Output probabilities necessarily sum to 1

# Softmax vs 1-vs-all

- Use softmax when we have mutually exclusive classes
  - Exclusive music genres: Pop, Rock, Jazz
- Use 1-vs-all when classes can overlap
  - A song can be Pop and Rock at the same time


# Preparing the data

- The data needs to be scaled and centered
- All columns need to be in the same range
- Height: 1.5 - 2.0
- Yearly income: 30000 - 200000
- Can use categorical data: SUNNY, CLOUDY, RAINY
- Need to be numerically encoded, introducing extra columns
- One-Hot encoding: SUNNY = {0, 0}; CLOUDY = {0, 1}; RAINY = {1, 0}
- Ordinal variables encoding is trickier: COLD, WARM, HOT
- We need to preserve the idea of ordering, but don't know the “distances” between COLD and WARM, WARM and HOT, etc.



# Overfitting


- Original  $f(x) = -0.7x + 0.5x^2 + 3$
- Gaussian error:  $\mathcal{N}(0, 0.7)$
- We fitted 4-degree polynomial
- And then 6-degree polynomial
- And then 9 degrees, which fits the data perfectly (9 points)
- However it will generalize poorly; it is **overfitting**
- The model is too complex; it is learning noise



# Bias/Variance Tradeoff

- Suppose we train the same model  $M$  on different training sets from the same population
- Bias -- a measure of the mean error of  $M$
- Variance -- a measure of how much the prediction of  $M$  differs from one training set to another
- High bias means underfitting; the model is too simple
- High variance means overfitting; the model is too complex
- In practice we can't eliminate both, hence the tradeoff

# Bias/Variance Tradeoff



# Bias/Variance Tradeoff

	<b><i>Bias</i></b>	<b><i>Variance</i></b>
<b><i>Low</i></b>	<ul style="list-style-type: none"><li>- Linear regression / linear data</li><li>- 3rd degree poly / quadratic data</li><li>- ANN with many nodes trained to completion</li></ul>	<ul style="list-style-type: none"><li>- Constant function</li><li>- Linear regression / quadratic data</li></ul>
<b><i>High</i></b>	<ul style="list-style-type: none"><li>- Constant function</li><li>- Linear regression / quadratic data</li><li>- ANN with few nodes applied to nonlinear data</li></ul>	<ul style="list-style-type: none"><li>- High degree poly</li><li>- ANN with many nodes trained to completion</li></ul>

# How to deal with overfitting


- More data; cancels the effects of noise
- Early stopping: stop before starting to overfit
- Model selection: select a model that overfits less
  - Cross-validation
- Dropout (NN): randomly drop unit contribution
  - Forces the model to learn with less input, hence less noise and co-adaptation
- Regularization: force a simpler model

# L2 Regularization

- Also known as Ridge or Tikhonov regularization
- Complex models have  $\mathbf{w}$  with a higher norm
- Change the loss function to force a lower norm
- Minimize:  $\mathcal{L}(w) = -\log L(w) + \lambda * \|w\|^2$
- High values for  $\mathbf{w}$  penalize the loss
- How do we choose  $\lambda$ ?
- In general, how do we choose the parameters of a model?

# Parameter tuning

- In general, *training set* error is not a reliable estimation of a model's performance
- We need a *test set*, which is never seen during training; used to report a model's performance
- If we need to tune parameters (e.g. regularization  $\lambda$ ), we set aside a *validation set*
- We train on training set, compare performance on different parameters on the validation set



# Cross-validation

- A more robust validation technique
- Instead of using a fixed training/validation split
- Perform multiple splits  $(t_1, v_1)$ ,  $(t_2, v_2)$ ,  $(t_3, v_3)$ , ... from the same training set
- Train using  $t_i$  and validate using  $v_i$
- The final error is the mean
- Multiple approaches:
  - Leave-one-out: validation set has size 1
  - K-fold: partition in K subsets; use K-1 for training and 1 for validation
    - Repeat K times
  - Repeated random split



# Demo

Q&A

